


Preprocessing as applied in training (patient M_DC-14-10926; exact torchvision transforms; last column: red / green / blue channel values after normalisation)



Columns 3-7 show each step at a fixed setting (rotation at its +15° limit, colour jitter at both limits); in training each is drawn at random, as in column 8. ConvNeXt-T, Swin-T and MedMamba get a new draw every epoch; the frozen Phikon probe was trained on one cached draw per image.