



Mean rank over 12 LOMO runs (1 = best, right); a bar joins methods that are not significantly different (Nemenyi)


Macro-F1 (Friedman p = 0.085; Nemenyi: no pair separated; Wilcoxon-Holm: MagScan-Mamba no encoder differs)




ECE (Friedman p = 0.0004; Nemenyi: 2 pair(s) separated; Wilcoxon-Holm: MagScan-Mamba differs from Phikon, ConvNeXt-T, Swin-T, MedMamba at p < 0.05)




Brier (Friedman p = 0.0077; Nemenyi: 2 pair(s) separated; Wilcoxon-Holm: MagScan-Mamba no encoder differs)



Risk @ 80% coverage (Friedman p = 0.043; Nemenyi: no pair separated; Wilcoxon-Holm: MagScan-Mamba no encoder differs)



E-AURC (Friedman p = 0.046; Nemenyi: 1 pair(s) separated; Wilcoxon-Holm: MagScan-Mamba differs from ConvNeXt-T, MedMamba at p < 0.05)



★ MagScan-Mamba ◆ Phikon (frozen FM) ● ConvNeXt-T ■ Swin-T ▲ MedMamba